BENCHMARKS:

Eigen: .0513
Boost : .107
STL: 1.44

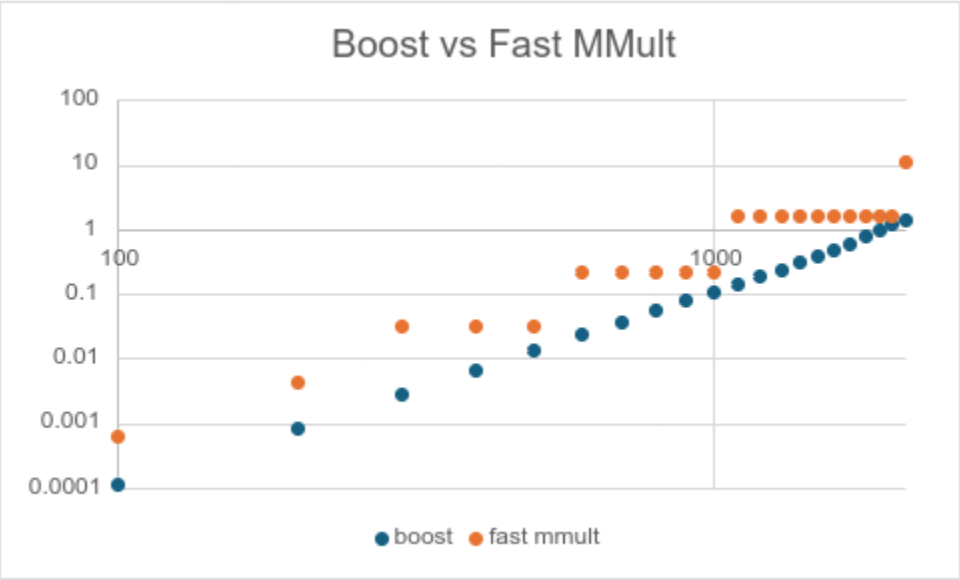

Boost seems to perform pretty well. One thing to note is the performance scales smoothly, and there are no sudden jumps in performance like mine.

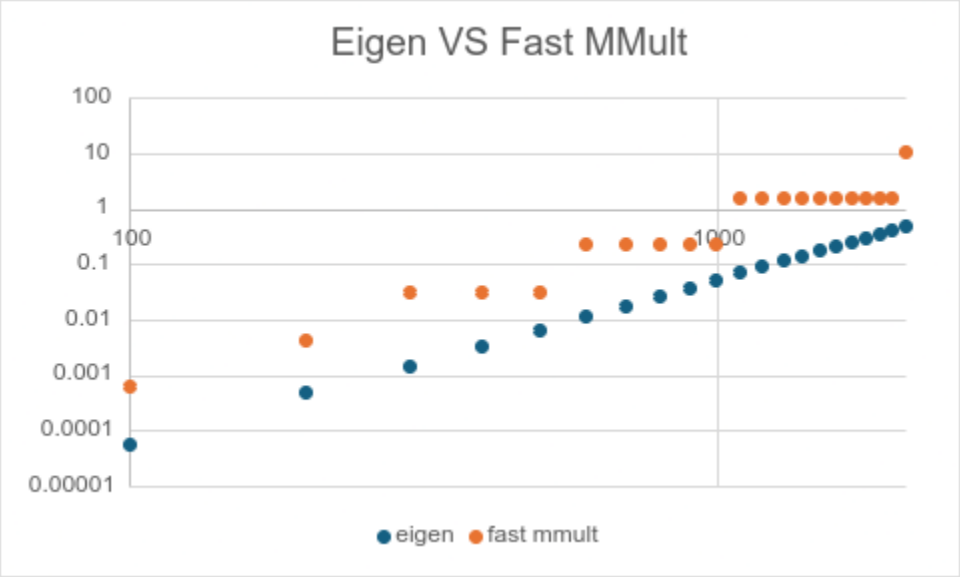

Eigen performs the best by far, and there is no scenario where my algorithm outperforms it. For a single core implementation, this is probably as good as it is going to get.

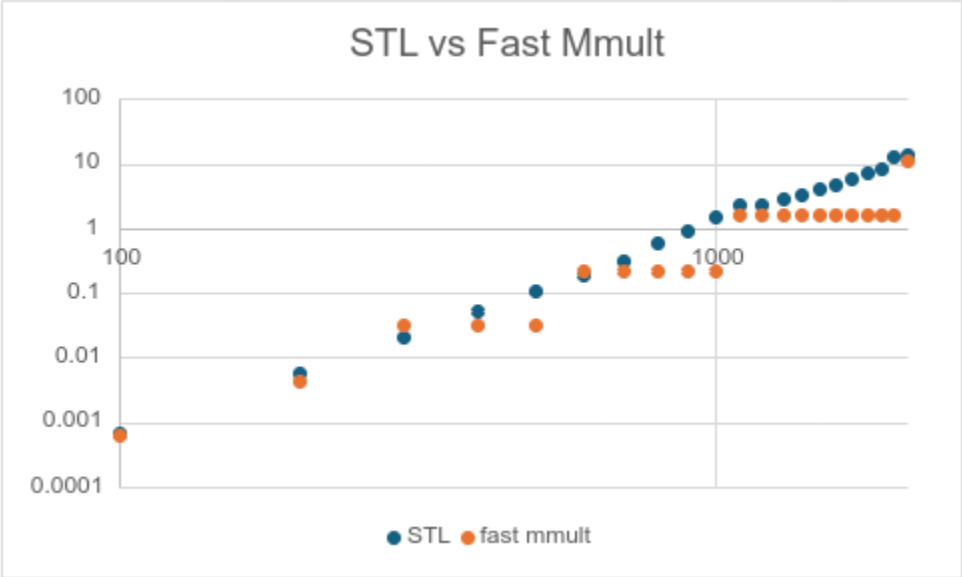

Stl just seems slower than my other testS. This is likely due to the overhead involved with handling vectors. It is much slower then my fast mmult implementation, and still slow compared to just a simple naive method. I think the reason this plot is a bit jumpy is that some other people were running jobs on nedc_130 while this was working, which gave me inconsistent results. I reran this job enough times, so I will just leave it for now.